

Group 18

M1 – SA: YOLOv5 Modeling**I. Libraries and Dependencies Installation**

Before the Roboflow dataset, that was produced in the previous weeks, is used to create YOLOv5 models, certain files and environments were prepared before testing the data. First, we cloned the *YOLOv5 from Ultralytics YOLOv5 Repository* and installed its dependencies. Since the dataset that will be used in this model will come from Roboflow, we also installed Roboflow's designated library, as seen in the code snippet below.

```
#clone forked YOLOv5 from Ultralytics YOLOv5 Repository
!git clone https://github.com/ultralytics/yolov5 # clone repo
%cd yolov5
%pip install -qr requirements.txt # install dependencies
%pip install -q roboflow # install roboflow

import torch
import os
from IPython.display import Image, clear_output # to display images

print(f"Setup complete. Using torch {torch.__version__} ({torch.cuda.get_device_properties(0).name if torch.cuda.is_available() else 'CPU'})")
```

The figure below displays the results of the above code snippet.

```
Cloning into 'yolov5'...
remote: Enumerating objects: 14308, done.
remote: Counting objects: 100% (66/66), done.
remote: Compressing objects: 100% (43/43), done.
remote: Total 14308 (delta 32), reused 49 (delta 23), pack-reused 14242
Receiving objects: 100% (14308/14308), 13.63 MiB | 17.40 MiB/s, done.
Resolving deltas: 100% (9832/9832), done.
/content/yolov5
| 182 kB 27.2 MB/s
| 62 kB 522 kB/s
| 1.6 MB 58.0 MB/s
| 42 kB 912 kB/s
| 145 kB 45.3 MB/s
| 138 kB 74.1 MB/s
| 54 kB 3.2 MB/s
| 178 kB 70.2 MB/s
| 67 kB 6.8 MB/s
| 62 kB 1.6 MB/s
Building wheel for wget (setup.py) ... done
Setup complete. Using torch 1.12.1+cu113 (Tesla T4)
```

Once the Roboflow dependency had been successfully installed, we imported Roboflow by applying the following code snippet:

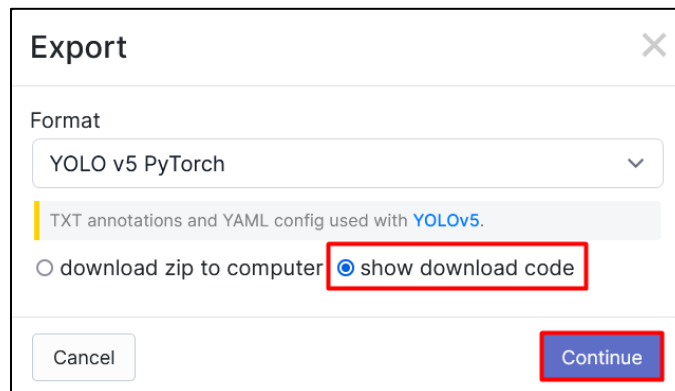
```
from roboflow import Roboflow #import roboflow for our datasets
rf = Roboflow(model_format="yolov5", notebook="ultralytics")
```

This code snippet will produce a result that will lead the programmer to a link where they can get an API key from Roboflow.

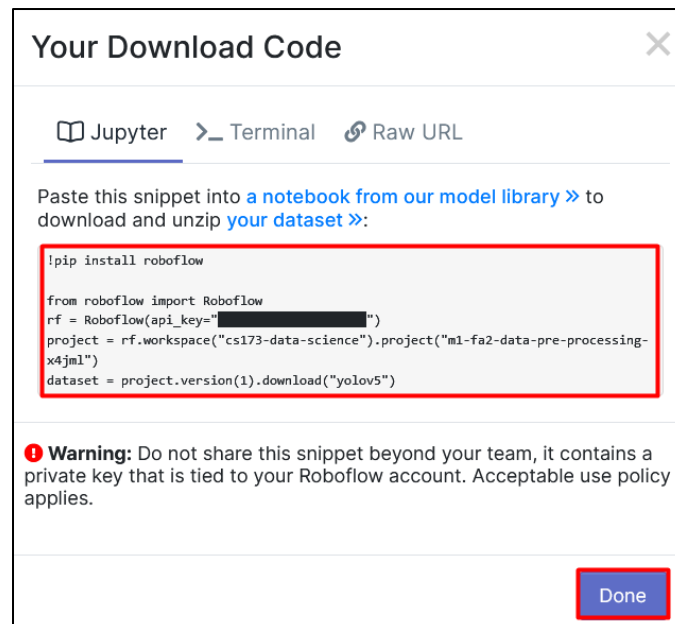
```
✓ [3] from roboflow import Roboflow #import roboflow for our datasets
0s rf = Roboflow(model_format="yolov5", notebook="ultralytics")
```

upload and label your dataset, and get an API KEY here: <https://app.roboflow.com/?model=yolov5&ref=ultralytics>

After getting the API key, you must go to your project's dataset and click on *Export*. Then, make sure that "show download code" is selected and click on *Continue*.

A dialog box titled "Export" with a close button (X) in the top right corner. It contains a "Format" dropdown menu set to "YOLO v5 PyTorch". Below the dropdown, there is a note: "TXT annotations and YAML config used with YOLOv5." At the bottom, there are two radio buttons: "download zip to computer" and "show download code", with the latter being selected and highlighted by a red box. At the very bottom, there are "Cancel" and "Continue" buttons, with the "Continue" button highlighted by a red box.

Copy the code snippet being shown below and click on Done.

A dialog box titled "Your Download Code" with a close button (X) in the top right corner. It has three tabs: "Jupyter" (selected), "Terminal", and "Raw URL". Below the tabs, there is a text prompt: "Paste this snippet into a notebook from our model library >> to download and unzip your dataset >>:". Below this prompt is a code snippet box with a red border containing the following code:

```
!pip install roboflow

from roboflow import Roboflow
rf = Roboflow(api_key="[REDACTED]")
project = rf.workspace("cs173-data-science").project("m1-fa2-data-pre-processing-x4jml")
dataset = project.version(1).download("yolov5")
```

Below the code snippet, there is a warning message: "Warning: Do not share this snippet beyond your team, it contains a private key that is tied to your Roboflow account. Acceptable use policy applies." At the bottom right, there is a "Done" button highlighted with a red box.

Back to the Google Colab notebook, we first set up the environment before pasting the code snippet we just copied.

```
# set up environment
os.environ["DATASET_DIRECTORY"] = "/content/datasets"
```

Then, paste the code snippet and run it. We wait for the Roboflow dataset to be downloaded.

```
!pip install roboflow

from roboflow import Roboflow
rf = Roboflow(api_key="RU6eGLCNP3Q4UwQqKWha")
project = rf.workspace("cs173-data-science").project("m1-fa2-data-pre-processing-x4jml")
dataset = project.version(1).download("yolov5")
```

II. Training the models (small, medium, and large)

After all necessary files and dependencies have been installed, we trained the model by implementing the *train.py* from the Ultralytics repository. The hyperparameters that were set for this model are as follows:

1. Image size = 416 pixels
2. Batch size = 16
3. Epochs = 50
4. YOLOv5 Version = YOLOv5s (Small)

Note: The “s” after “YOLOv5” indicates the version scale (size) of the model. One requirement of this project is to create three YOLOv5 models for scales small, medium, and large; therefore, we adjusted the YOLOv5 versions by changing the last letter to the first letter of their appropriate version scale. Since the process for each model will be the same for each model of a different weight, the model with version YOLOv5s will only be discussed in this section.

To apply the ideal hyperparameters mentioned above, the following code snippet was run, yielding the results from the succeeding figures:

```
!python train.py --img 416 --batch 16 --epochs 50 --
data {dataset.location}/data.yaml --weights yolov5s.pt --
name yolov5s_results -cache
```

train: weights=yolov5s.pt, cfg=, data=/content/datasets/M1---FA2-Data-Pre-Processing-1/data.yaml, hyp=data/hyps/hyp.scratch-low.yaml, epochs=50, batch_size=16, img
github: up to date with <https://github.com/ultralytics/yolov5>
YOLOv5 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

hyperparameters: lr=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=0.05, cls=0.5, cls_pw=1.0
ClearML: run 'pip install clearml' to automatically track, visualize and remotely train YOLOv5 in ClearML

Comet: run 'pip install comet_ml' to automatically track and visualize YOLOv5 runs in Comet

TensorBoard: Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>

Downloading <https://ultralytics.com/assets/Arial.ttf> to /root/.config/Ultralytics/Arial.ttf...

100% 755k/755k [00:00<00:00, 176MB/s]

Downloading <https://github.com/ultralytics/yolov5/releases/download/v7.0/yolov5s.pt> to yolov5s.pt...

100% 14.1M/14.1M [00:00<00:00, 304MB/s]

Overriding model.yaml nc=80 with nc=5

		from	n	params	module	arguments
0		-1	1	3520	models.common.Conv	[3, 32, 6, 2, 2]
1		-1	1	18560	models.common.Conv	[32, 64, 3, 2]
2		-1	1	18816	models.common.C3	[64, 64, 1]
3		-1	1	73984	models.common.Conv	[64, 128, 3, 2]
4		-1	2	115712	models.common.C3	[128, 128, 2]
5		-1	1	295424	models.common.Conv	[128, 256, 3, 2]
6		-1	3	625152	models.common.C3	[256, 256, 3]
7		-1	1	1180672	models.common.Conv	[256, 512, 3, 2]
8		-1	1	1182720	models.common.C3	[512, 512, 1]
9		-1	1	656896	models.common.SPPF	[512, 512, 5]
10		-1	1	131584	models.common.Conv	[512, 256, 1, 1]
11		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12		[-1, 6]	1	0	models.common.Concat	[1]
13		-1	1	361984	models.common.C3	[512, 256, 1, False]
14		-1	1	33024	models.common.Conv	[256, 128, 1, 1]
15		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16		[-1, 4]	1	0	models.common.Concat	[1]
17		-1	1	90880	models.common.C3	[256, 128, 1, False]
18		-1	1	147712	models.common.Conv	[128, 128, 3, 2]
19		[-1, 14]	1	0	models.common.Concat	[1]
20		-1	1	296448	models.common.C3	[256, 256, 1, False]

21		-1	1	590336	models.common.Conv	[256, 256, 3, 2]
22		[-1, 10]	1	0	models.common.Concat	[1]
23		-1	1	1182720	models.common.C3	[512, 512, 1, False]
24		[17, 20, 23]	1	26970	models.yolo.Detect	[5, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90, 156, 198, 373, 326]]]

Model summary: 214 layers, 7033114 parameters, 7033114 gradients, 16.0 GFLOPs

Transferred 343/349 items from yolov5s.pt

AMP: checks passed

optimizer: SGD(lr=0.01) with parameter groups 57 weight(decay=0.0), 60 weight(decay=0.0005), 60 bias

augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01), CLAHE(p=0.01, clip_limit=(1, 4.0), tile_grid_size=(8, 8))

train: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/train/labels... 1799 images, 123 backgrounds, 0 corrupt: 100% 1799/1799 [00:00<00:00, 1984.61it/s]

train: New cache created: /content/datasets/M1---FA2-Data-Pre-Processing-1/train/labels.cache

train: Caching images (0.96B ram): 100% 1799/1799 [00:10<00:00, 175.50it/s]

val: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels... 180 images, 17 backgrounds, 0 corrupt: 100% 180/180 [00:00<00:00, 609.53it/s]

val: New cache created: /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels.cache

val: Caching images (0.16B ram): 100% 180/180 [00:01<00:00, 101.44it/s]

AutoAnchor: 4.75 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset

Plotting labels to runs/train/yolov5s_results/labels.jpg...

Image sizes 416 train, 416 val

Using 2 dataloader workers

Logging results to runs/train/yolov5s_results

Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
0/49	1.71G	0.09907	0.03477	0.03224	40	416: 100% 113/113 [00:23<00:00, 4.72it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:02<00:00, 2.29it/s]
	all	180	802	0.825	0.0712	0.0164 0.00435
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
1/49	2.07G	0.07245	0.02921	0.01683	52	416: 100% 113/113 [00:19<00:00, 5.83it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 5.32it/s]
	all	180	802	0.686	0.0944	0.059 0.0217
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
2/49	2.07G	0.06545	0.0249	0.01445	49	416: 100% 113/113 [00:19<00:00, 5.93it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 5.57it/s]
	all	180	802	0.478	0.135	0.0737 0.0298

	all	180	802	0.465	0.213	0.215	0.109	
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
46/49	2.07G	0.02083	0.01162	0.00762	40	416: 100% 113/113 [00:20<00:00, 5.63it/s]		
Class	Images	Instances	P	R	mAP50	mAP50-95: 100% 6/6 [00:00<00:00, 6.07it/s]		
all	180	802	0.47	0.204	0.209	0.103		
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
47/49	2.07G	0.02044	0.01096	0.00753	35	416: 100% 113/113 [00:18<00:00, 6.18it/s]		
Class	Images	Instances	P	R	mAP50	mAP50-95: 100% 6/6 [00:01<00:00, 5.90it/s]		
all	180	802	0.462	0.21	0.216	0.108		
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
48/49	2.07G	0.02077	0.0117	0.00738	45	416: 100% 113/113 [00:18<00:00, 6.02it/s]		
Class	Images	Instances	P	R	mAP50	mAP50-95: 100% 6/6 [00:01<00:00, 5.79it/s]		
all	180	802	0.462	0.22	0.219	0.109		
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
49/49	2.07G	0.02038	0.01132	0.007088	41	416: 100% 113/113 [00:18<00:00, 6.08it/s]		
Class	Images	Instances	P	R	mAP50	mAP50-95: 100% 6/6 [00:01<00:00, 5.94it/s]		
all	180	802	0.457	0.215	0.218	0.109		
50 epochs completed in 0.283 hours.								
Optimizer stripped from runs/train/yolov5s_results/weights/last.pt, 14.3MB								
Optimizer stripped from runs/train/yolov5s_results/weights/best.pt, 14.3MB								
Validating runs/train/yolov5s_results/weights/best.pt...								
Fusing layers...								
Model summary: 157 layers, 7023610 parameters, 0 gradients, 15.8 GFLOPs								
Class	Images	Instances	P	R	mAP50	mAP50-95: 100% 6/6 [00:02<00:00, 2.27it/s]		
all	180	802	0.479	0.217	0.224	0.111		
Full-Faced	180	178	0.374	0.365	0.318	0.104		
Half-Faced	180	158	0.223	0.228	0.144	0.0485		
Invalid	180	27	0	0	0.0162	0.00572		
Not Wearing Helmet	180	29	1	0	0.0158	0.00875		
Rider	180	410	0.796	0.494	0.627	0.389		
Results saved to runs/train/yolov5s_results								

Below are the codes and results for the **medium-scaled** model.

```
!python train.py --img 416 --batch 16 --epochs 50 --
data {dataset.location}/data.yaml --weights yolov5m.pt --
name yolov5m_results --cache ram
```

train: weights=yolov5m.pt, cfg=, data=/content/datasets/M1---FA2-Data-Pre-Processing-1/data.yaml, hyp=data/hyps/hyp.scratch-low.yaml, epochs=5
github: up to date with <https://github.com/ultralytics/yolov5> ✓
YOLOv5 🚀 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

hyperparameters: lr0=0.01, lr1=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=0.05
ClearML: run 'pip install clearml' to automatically track, visualize and remotely train YOLOv5 🚀 in ClearML
Comet: run 'pip install comet_ml' to automatically track and visualize YOLOv5 🚀 runs in Comet
TensorBoard: Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>
Downloading <https://ultralytics.com/assets/Arial.ttf> to /root/.config/ultralytics/Arial.ttf...
100% 755k/755k [00:00<00:00, 93.3MB/s]
Downloading <https://github.com/ultralytics/yolov5/releases/download/v7.0/yolov5m.pt> to yolov5m.pt...
100% 40.8M/40.8M [00:05<00:00, 7.44MB/s]

Overriding model.yaml nc=80 with nc=5

	from	n	params	module	arguments
0	-1	1	5280	models.common.Conv	[3, 48, 6, 2, 2]
1	-1	1	41664	models.common.Conv	[48, 96, 3, 2]
2	-1	2	65280	models.common.C3	[96, 96, 2]
3	-1	1	166272	models.common.Conv	[96, 192, 3, 2]
4	-1	4	444672	models.common.C3	[192, 192, 4]
5	-1	1	664320	models.common.Conv	[192, 384, 3, 2]
6	-1	6	2512896	models.common.C3	[384, 384, 6]
7	-1	1	2655744	models.common.Conv	[384, 768, 3, 2]
8	-1	2	4134912	models.common.C3	[768, 768, 2]
9	-1	1	1476864	models.common.SPPF	[768, 768, 5]
10	-1	1	295680	models.common.Conv	[768, 384, 1, 1]
11	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[-1, 6]	1	0	models.common.Concat	[1]

```

13         -1 2 1182720 models.common.C3 [768, 384, 2, False]
14         -1 1 74112 models.common.Conv [384, 192, 1, 1]
15         -1 1 0 torch.nn.modules.upsampling.Upsample [None, 2, 'nearest']
16     [-1, 4] 1 0 models.common.Concat [1]
17         -1 2 296448 models.common.C3 [384, 192, 2, False]
18         -1 1 332160 models.common.Conv [192, 192, 3, 2]
19     [-1, 14] 1 0 models.common.Concat [1]
20         -1 2 1035264 models.common.C3 [384, 384, 2, False]
21         -1 1 1327872 models.common.Conv [384, 384, 3, 2]
22     [-1, 10] 1 0 models.common.Concat [1]
23         -1 2 4134912 models.common.C3 [768, 768, 2, False]
24     [17, 20, 23] 1 40410 models.yolo.Detect [5, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 90
Model summary: 291 layers, 20887482 parameters, 20887482 gradients, 48.3 GFLOPs

```

Transferred 475/481 items from yolov5m.pt

AMP: checks passed ☒

optimizer: SGD(lr=0.01) with parameter groups 79 weight(decay=0.0), 82 weight(decay=0.0005), 82 bias

augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01), CLAHE(p=0.01, clip_limit=(1, 4.0), til

train: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/train/labels... 1799 images, 123 backgrounds, 0 corrupt: 100% 1799/1799 [00:00:00:00]

train: New cache created: /content/datasets/M1---FA2-Data-Pre-Processing-1/train/labels.cache

train: Caching images (0.9GB ram): 100% 1799/1799 [00:10:00:00, 176.03it/s]

val: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels... 180 images, 17 backgrounds, 0 corrupt: 100% 180/180 [00:00:00:00]

val: New cache created: /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels.cache

val: Caching images (0.1GB ram): 100% 180/180 [00:01:00:00, 106.00it/s]

AutoAnchor: 4.75 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset ☒

Plotting labels to runs/train/yolov5m_results/labels.jpg...

Image sizes 416 train, 416 val

Using 2 dataloader workers

Logging results to runs/train/yolov5m_results

Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
0/49	2.93G	0.09309	0.03582	0.03039	40	416: 100% 113/113 [00:30<00:00, 3.69it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:02<00:00, 2.14it/s]
	all	180	802	0.847	0.0532	0.0327 0.013
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
1/49	3.47G	0.06917	0.02783	0.01633	52	416: 100% 113/113 [00:26<00:00, 4.28it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.17it/s]
	all	180	802	0.679	0.0914	0.0661 0.0252
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
2/49	3.47G	0.0632	0.02317	0.01427	49	416: 100% 113/113 [00:24<00:00, 4.57it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.23it/s]
	all	180	802	0.48	0.14	0.0795 0.0344
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
3/49	3.47G	0.05482	0.02079	0.01267	57	416: 100% 113/113 [00:24<00:00, 4.60it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.20it/s]
	all	180	802	0.573	0.168	0.149 0.0678
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
4/49	3.47G	0.04713	0.01881	0.01179	38	416: 100% 113/113 [00:24<00:00, 4.58it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.18it/s]
	all	180	802	0.629	0.226	0.192 0.0897
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
47/49	3.47G	0.01714	0.008929	0.005369	35	416: 100% 113/113 [00:24<00:00, 4.55it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.22it/s]
	all	180	802	0.498	0.267	0.253 0.125
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
48/49	3.47G	0.0174	0.009438	0.005209	45	416: 100% 113/113 [00:24<00:00, 4.54it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.23it/s]
	all	180	802	0.494	0.262	0.251 0.126
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
49/49	3.47G	0.01705	0.009126	0.005106	41	416: 100% 113/113 [00:24<00:00, 4.53it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 4.26it/s]
	all	180	802	0.303	0.264	0.256 0.13

50 epochs completed in 0.381 hours.

Optimizer stripped from runs/train/yolov5m_results/weights/last.pt, 42.1MB

Optimizer stripped from runs/train/yolov5m_results/weights/best.pt, 42.1MB

Validating runs/train/yolov5m_results/weights/best.pt...

Fusing layers...

Model summary: 212 layers, 20869098 parameters, 0 gradients, 47.9 GFLOPs

	Class	Images	Instances	P	R	mAP50	mAP50-95: 100% 6/6 [00:04<00:00, 1.45it/s]
	all	180	802	0.303	0.264	0.256	0.13
	Full-Faced	180	178	0.398	0.433	0.411	0.16
	Half-Faced	180	158	0.261	0.272	0.193	0.0734
	Invalid	180	27	0.0592	0.111	0.0193	0.0103
	Not Wearing Helmet	180	29	0	0	0.0188	0.0086
	Rider	180	410	0.799	0.503	0.637	0.398

Results saved to runs/train/yolov5m_results

Lastly, these are the codes and results for the **large-scaled** model.

```
!python train.py --img 416 --batch 16 --epochs 50 --
data {dataset.location}/data.yaml --weights yolov5l.pt --
name yolov5l_results --cache ram
```

train: weights=yolov5l.pt, cfg=, data=/content/datasets/M1---FA2-Data-Pre-Processing-1/data.yaml, hyp=data/hyps/hyp.scratch-low.yaml, epochs=5
github: up to date with <https://github.com/ultralytics/yolov5> ✓
YOLOv5 🚀 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

hyperparameters: lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=0.05
ClearML: run 'pip install clearml' to automatically track, visualize and remotely train YOLOv5 🚀 in ClearML
Comet: run 'pip install comet_ml' to automatically track and visualize YOLOv5 🚀 runs in Comet
TensorBoard: Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>
Overriding Model.yaml nc=80 with nc=5

		from	n	params	module	arguments
0		-1	1	7040	models.common.Conv	[3, 64, 6, 2, 2]
1		-1	1	73984	models.common.Conv	[64, 128, 3, 2]
2		-1	3	156928	models.common.C3	[128, 128, 3]
3		-1	1	295424	models.common.Conv	[128, 256, 3, 2]
4		-1	6	1118208	models.common.C3	[256, 256, 6]
5		-1	1	1180672	models.common.Conv	[256, 512, 3, 2]
6		-1	9	6433792	models.common.C3	[512, 512, 9]
7		-1	1	4720640	models.common.Conv	[512, 1024, 3, 2]
8		-1	3	9971712	models.common.C3	[1024, 1024, 3]
9		-1	1	2624512	models.common.SPPF	[1024, 1024, 5]
10		-1	1	525312	models.common.Conv	[1024, 512, 1, 1]
11		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[-1, 6]	1	0	models.common.Concat	[1]	
13		-1	3	2757632	models.common.C3	[1024, 512, 3, False]
14		-1	1	131584	models.common.Conv	[512, 256, 1, 1]
15		-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16	[-1, 4]	1	0	models.common.Concat	[1]	
17		-1	3	690688	models.common.C3	[512, 256, 3, False]
18		-1	1	590336	models.common.Conv	[256, 256, 3, 2]
19	[-1, 14]	1	0	models.common.Concat	[1]	
20		-1	3	2495488	models.common.C3	[512, 512, 3, False]
21		-1	1	2360320	models.common.Conv	[512, 512, 3, 2]
22	[-1, 10]	1	0	models.common.Concat	[1]	
23		-1	3	9971712	models.common.C3	[1024, 1024, 3, False]
24	[17, 20, 23]	1	53850	models.yolo.Detect	[5, [[10, 13, 16, 30, 33, 23], [30, 61, 62, 45, 59, 119], [116, 96	

Model summary: 368 layers, 46159834 parameters, 46159834 gradients, 108.3 GFLOPs

Transferred 607/613 items from yolov5l.pt

AMP: checks passed ✓

optimizer: SGD(lr=0.01) with parameter groups 101 weight(decay=0.0), 104 weight(decay=0.0005), 104 bias

augmentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01), CLAHE(p=0.01, clip_limit=(1, 4.0), til

train: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/train/labels.cache... 1799 images, 123 backgrounds, 0 corrupt: 100% 1799/1795

train: Caching images (0.9GB ram): 100% 1799/1799 [00:11<00:00, 161.85it/s]

val: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels.cache... 180 images, 17 backgrounds, 0 corrupt: 100% 180/180 [00:0

val: Caching images (0.1GB ram): 100% 180/180 [00:02<00:00, 84.86it/s]

AutoAnchor: 4.75 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset ✓

Plotting labels to runs/train/yolov5l_results2/labels.jpg...

Image sizes 416 train, 416 val

Using 2 dataloader workers

Logging results to runs/train/yolov5l_results2

Starting training for 50 epochs...

Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
0/49	4.77G	0.09032	0.03636	0.03056	40	416: 100% 113/113 [00:39<00:00, 2.88it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:02<00:00, 2.86it/s]
	all	180	802	0.889	0.0671	0.0629 0.0243

Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size
1/49	5.49G	0.06823	0.0269	0.01615	52	416: 100% 113/113 [00:34<00:00, 3.30it/s]
	Class	Images	Instances	P	R	mAP50 mAP50-95: 100% 6/6 [00:01<00:00, 3.39it/s]
	all	180	802	0.864	0.0751	0.0728 0.0322

Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
46/49	5.49G	0.01627	0.008554	0.004425	40	416: 100% 113/113 [00:33<00:00, 3.40it/s]		
	Class	Images	Instances	P	R	mAP50	mAP50-95:	100% 6/6 [00:01<00:00, 3.47it/s]
	all	180	802	0.335	0.259	0.261	0.134	
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
47/49	5.49G	0.01575	0.007904	0.004274	35	416: 100% 113/113 [00:33<00:00, 3.40it/s]		
	Class	Images	Instances	P	R	mAP50	mAP50-95:	100% 6/6 [00:01<00:00, 3.54it/s]
	all	180	802	0.35	0.262	0.268	0.136	
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
48/49	5.49G	0.01597	0.008478	0.004064	45	416: 100% 113/113 [00:33<00:00, 3.41it/s]		
	Class	Images	Instances	P	R	mAP50	mAP50-95:	100% 6/6 [00:01<00:00, 3.56it/s]
	all	180	802	0.376	0.277	0.276	0.141	
Epoch	GPU_mem	box_loss	obj_loss	cls_loss	Instances	Size		
49/49	5.49G	0.01552	0.008144	0.003962	41	416: 100% 113/113 [00:33<00:00, 3.40it/s]		
	Class	Images	Instances	P	R	mAP50	mAP50-95:	100% 6/6 [00:01<00:00, 3.54it/s]
	all	180	802	0.361	0.254	0.271	0.142	

50 epochs completed in 0.515 hours.

Optimizer stripped from runs/train/yolov5l_results2/weights/last.pt, 92.8MB

Optimizer stripped from runs/train/yolov5l_results2/weights/best.pt, 92.8MB

Validating runs/train/yolov5l_results2/weights/best.pt...

Fusing layers...

Model summary: 267 layers, 46129818 parameters, 0 gradients, 107.7 GFLOPs

Class	Images	Instances	P	R	mAP50	mAP50-95:	100% 6/6 [00:03<00:00, 1.71it/s]
all	180	802	0.328	0.263	0.271	0.143	
Full-Faced	180	178	0.432	0.452	0.404	0.162	
Half-Faced	180	158	0.33	0.297	0.25	0.108	
Invalid	180	27	0.0311	0.037	0.02	0.00761	
Not Wearing Helmet	180	29	0	0	0.0283	0.00866	
Rider	180	410	0.846	0.527	0.651	0.428	

Results saved to runs/train/yolov5l_results2

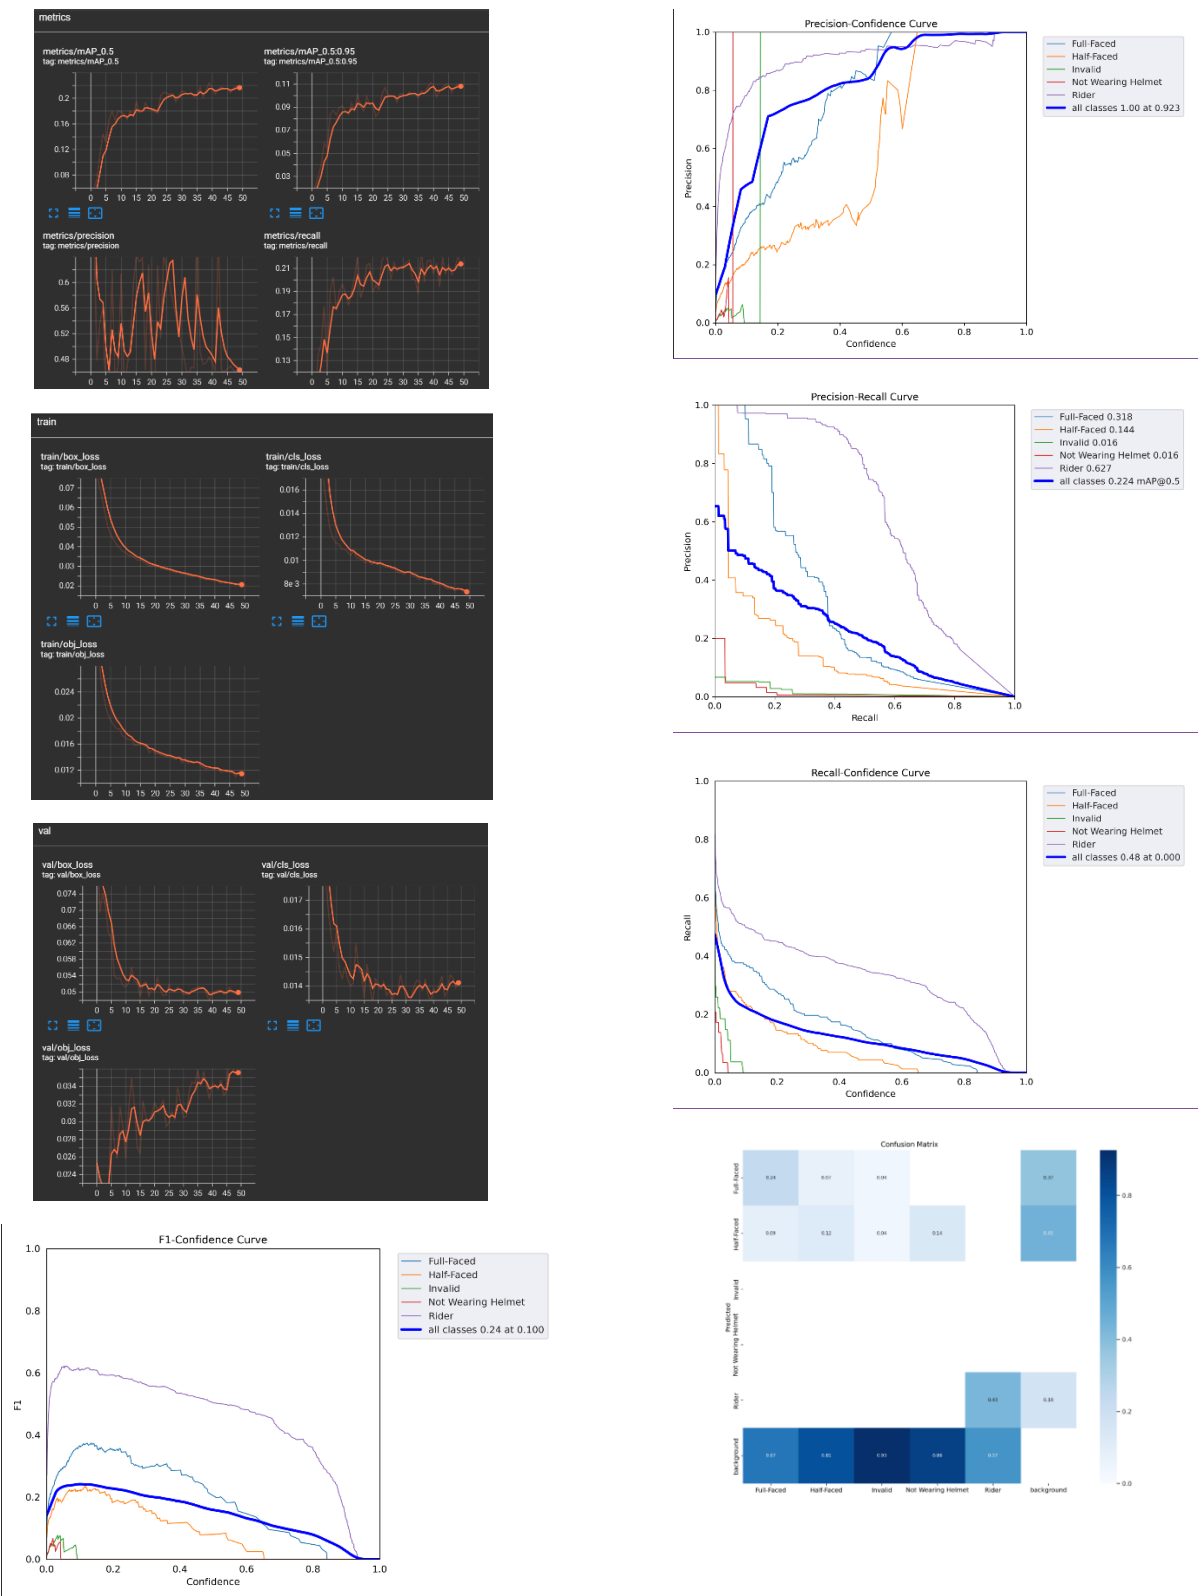
III. TensorBoard

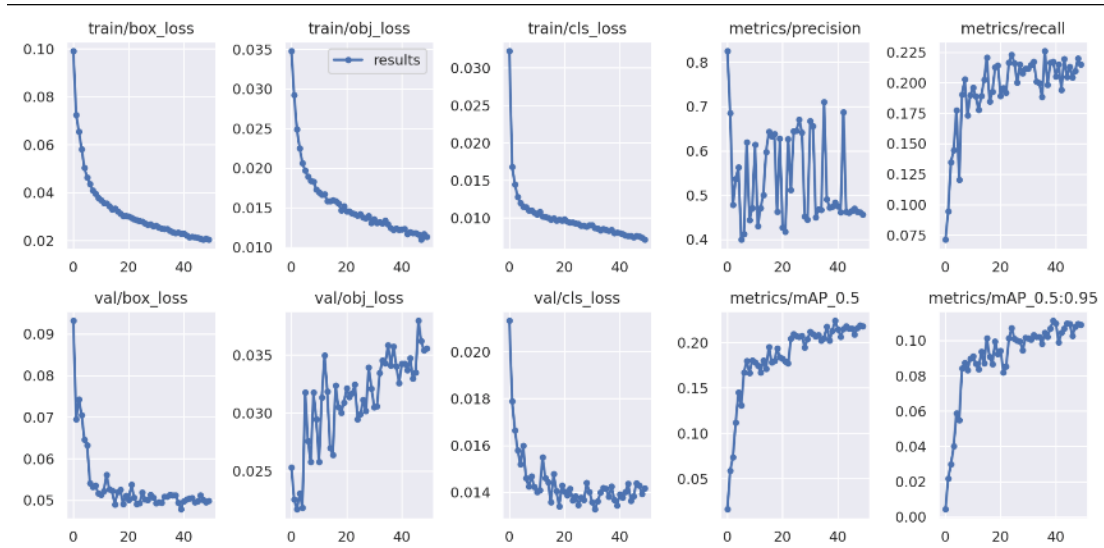
In this project, TensorBoard was used to portray the results of the training process through scalar metrics, images, graphs, and time series. A few examples of well-known scalar metrics are *mAP* (mean average precision), *precision*, and *recall*. These values determine whether the trained model is underfitting or overfitting. Under images and graphs, one will see the performance of the model through graphs that represent the *F1 curve*, *PC curve*, *confusion matrix*, and the model's overall results. The time series shows all runtimes of the model's training process.

To launch TensorBoard, the following code snippet was run:


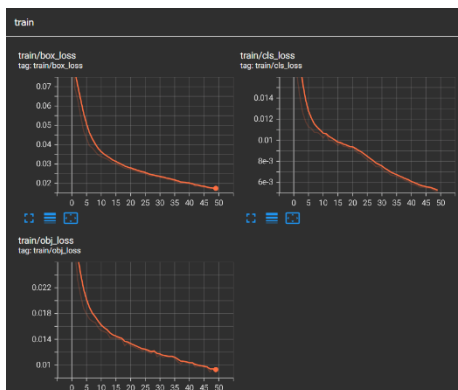
```
%load_ext tensorboard
%tensorboard --logdir runs
```

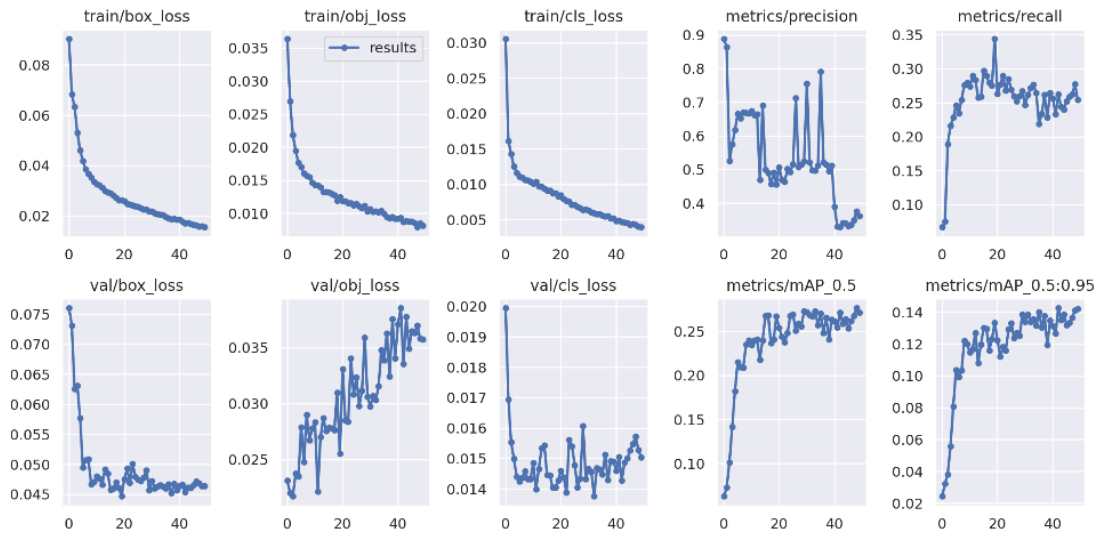





The graphical results of the YOLOv5s (small) model are seen in the figures below.






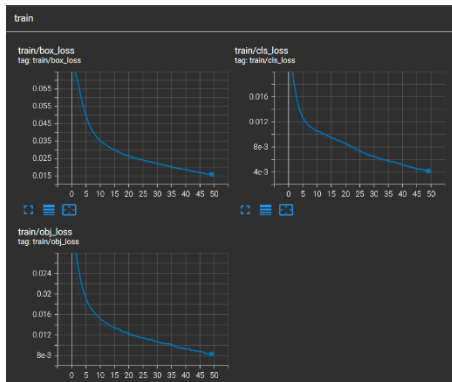

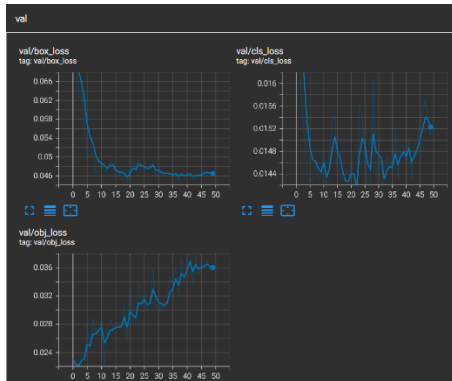
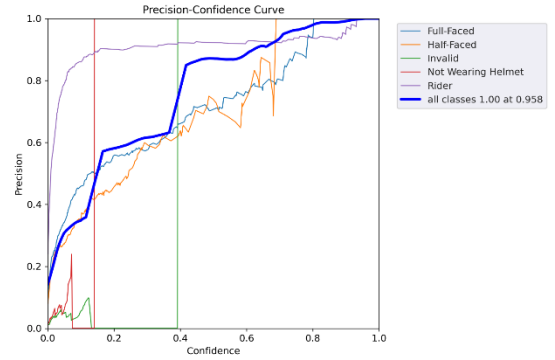
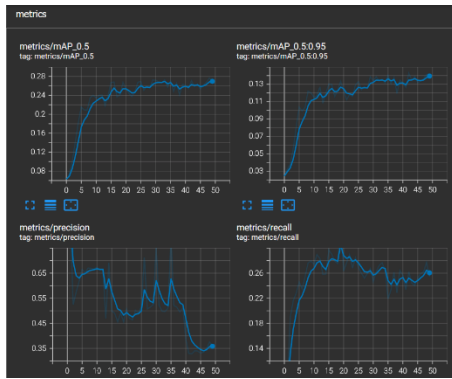



The results for the YOLOv5s (**medium**) model are seen below:





Finally, the results for the YOLOv5s (**large**) model are in the following figures:





IV. Testing the models (small, medium, and large)

After training the model using the training dataset, the testing dataset, which consists of 4,500 frames from traffic surveillance footage, will now be used to test the model. The purpose of this dataset is to measure the performance metrics and detection capabilities of the model. To fulfill this purpose, *detect.py* from the Ultralytics repository was utilized to detect objects and predict where the bounding boxes from the test images are to be placed.


```
!python detect.py --weights runs/train/yolov5s_results/weights/best.pt --img 416 --conf 0.1 --source {dataset.location}/test/images
```

Output:

```
#detecting and predicting using the test dataset
python detect.py --weights runs/train/yolov5s_results/weights/best.pt --img 416 --conf 0.1 --source {dataset.location}/test/images

detect: weights='runs/train/yolov5s_results/weights/best.pt', source=/content/datasets/M1---FA2-Data-Pre-Processing-1/test/images, data=data/coco128.yaml, imgs=[
YOLOv5 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

Fusing layers...
Model summary: 157 layers, 7023610 parameters, 0 gradients, 15.8 GFLOPs
image 1/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-100.jpg.rf.dcfbacfe8d133a52a673353a0a4453d...
image 2/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-108.jpg.rf.f46b579b4ffe244f48820dc5139091b2...
image 3/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-112.jpg.rf.c4eca008ac8f91fb5c345f5e819ef415...
image 4/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-121.jpg.rf.36be41fd6e37a2a50952b0e4b784deec...
image 5/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-127.jpg.rf.6a608a9f9280a8abe4c024866ada407f...
image 6/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-144.jpg.rf.f4aa9ccf658a47b06eb7eb3501e2e018...
image 7/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-173.jpg.rf.76d564d59076ee0b8073bef1d17362c...
image 8/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-183.jpg.rf.2072b184227de0ea90766a0da38f293b...
image 9/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-18.jpg.rf.cabe20fb408c2c960d848cd9e9bd3bf...
image 10/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-191.jpg.rf.4483a68b586f14b6d2f1cf821a14220...
image 11/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-192.jpg.rf.f413019c3cfae8640cd14429d6e4f6e6...
image 12/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-193.jpg.rf.d18bacb9496daa2de8f533a6d0d8685...
image 13/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-200.jpg.rf.b182449eb1494201abc4b24d6b15bd7...
image 14/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-206.jpg.rf.7b8738ae9c034677a7971c341c8627a...
image 15/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-209.jpg.rf.746e98069aa492e43006c1a5525be809...
image 16/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-223.jpg.rf.c17cf723dd1f95c3b2b77ce1142716...
image 17/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-225.jpg.rf.b974cad46afef6a8945f0ff877204f480...
image 18/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-237.jpg.rf.b89be9b397b700be27c5ad415fed85...
image 19/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-249.jpg.rf.3f55266c55797eb7fd3180112391e153...
image 20/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-254.jpg.rf.06c38e54c8b9c2d13afced6817daca...
image 21/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-256.jpg.rf.39cc97a2ee129ec96054a70e4e5972a...
image 22/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-259.jpg.rf.4656bca09500f7baba3461ac556b0f...
image 23/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-264.jpg.rf.908b8999d11cd96825cd6a264ac22895...
image 24/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-284.jpg.rf.f1cbcdce41cd719eb230b3e94ef2aa5a...
image 25/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-291.jpg.rf.8e82c1694799e1e8787460e8196a7bbe...
image 26/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-293.jpg.rf.70a9f5bdfda5955f0c01f5f6ae9eb564...
image 27/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-35.jpg.rf.8cae35ccccc606f5346ee1668dc7469e...
image 28/90 /content/datasets/M1---FA2-Data-Pre-Processing-1/test/images/XVR_ch9_main_20221004170008_20221004180008_mp4-45.jpg.rf.a4475cd731119823723748695a5b680...
59s completed at 11:32 PM
```



To display all test images with the bounding boxes placed on their predictions, the following libraries were imported:

```
import glob
from IPython.display import Image, display
```

Then, to display the inferred test images on the next cell by retrieving the images from the specified directory, the code snippet below was run.

```
for imageName in glob.glob('/content/yolov5/runs/detect/exp2/*.jpg'):
    #assuming JPG
    display(Image(filename=imageName))
    print("\n")
```


The figures below are a few examples of the inferred (with bounding boxes) test images.





Next, the following code was run to produce the table output in the figure below. The results will be saved as a file named “best.pt.” The code snippet will produce a summary of the model’s values through a table, as its classes are divided among Full-Faced, Half-Faced, Invalid, Not Wearing Helmet, and Rider and classified according to the number of images, instances, precision (P), recall (R), mAP50, and mAP50-95.

```
!python val.py --weights './runs/train/yolov5s_results/weights/best.pt' --
data {dataset.location}/data.yaml --img 416
```

Output for the small-scaled model:

```
val: data=/content/datasets/M1---FA2-Data-Pre-Processing-1/data.yaml, weights=['./runs/train/yolov5s_results/weights/best.pt'], batch_size=32, imgsz=416,
YOLOv5 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

Fusing layers...
Model summary: 157 layers, 7023610 parameters, 0 gradients, 15.8 GFLOPs
val: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels.cache... 180 images, 17 backgrounds, 0 corrupt: 100% 180/180 [00:00<?, ?it/s]

```

Class	Images	Instances	P	R	mAP50	mAP50-95
all	180	802	0.485	0.218	0.226	0.111
Full-Faced	180	178	0.394	0.369	0.325	0.105
Half-Faced	180	158	0.229	0.228	0.144	0.0483
Invalid	180	27	0	0	0.0161	0.00568
Not Wearing Helmet	180	29	1	0	0.0157	0.00875
Rider	180	410	0.801	0.495	0.626	0.389

```
Speed: 0.3ms pre-process, 3.4ms inference, 2.3ms NMS per image at shape (32, 3, 416, 416)
Results saved to runs/val/exp3
```

Output for the medium-scaled model:

```
val: data=/content/datasets/M1---FA2-Data-Pre-Processing-1/data.yaml, weights=['./runs/train/yolov5m_results/weights/best.pt'], batch_size=32, i
YOLOv5 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

Fusing layers...
Model summary: 212 layers, 20869098 parameters, 0 gradients, 47.9 GFLOPs
val: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels.cache... 180 images, 17 backgrounds, 0 corrupt: 100% 180/180 [00:00<

```

Class	Images	Instances	P	R	mAP50	mAP50-95
all	180	802	0.303	0.264	0.256	0.13
Full-Faced	180	178	0.395	0.427	0.409	0.158
Half-Faced	180	158	0.263	0.272	0.194	0.0732
Invalid	180	27	0.0592	0.111	0.0191	0.01
Not Wearing Helmet	180	29	0	0	0.02	0.00879
Rider	180	410	0.797	0.507	0.637	0.398

```
Speed: 0.6ms pre-process, 6.9ms inference, 2.2ms NMS per image at shape (32, 3, 416, 416)
Results saved to runs/val/exp2
```

Output for the large-scaled model:

```
val: data=/content/datasets/M1---FA2-Data-Pre-Processing-1/data.yaml, weights=['./runs/train/yolov5l_results2/weights/best.pt'], batch_size=32,
YOLOv5 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

Fusing layers...
Model summary: 267 layers, 46129818 parameters, 0 gradients, 107.7 GFLOPs
val: Scanning /content/datasets/M1---FA2-Data-Pre-Processing-1/valid/labels.cache... 180 images, 17 backgrounds, 0 corrupt: 100% 180/180 [00:00<

```

Class	Images	Instances	P	R	mAP50	mAP50-95
all	180	802	0.328	0.262	0.268	0.142
Full-Faced	180	178	0.427	0.447	0.39	0.16
Half-Faced	180	158	0.333	0.297	0.25	0.108
Invalid	180	27	0.0311	0.037	0.02	0.00737
Not Wearing Helmet	180	29	0	0	0.028	0.00884
Rider	180	410	0.849	0.527	0.651	0.428

```
Speed: 0.5ms pre-process, 11.6ms inference, 2.3ms NMS per image at shape (32, 3, 416, 416)
Results saved to runs/val/exp
```

V. Discussion of results for each model

As demonstrated above, each model was tested and validated using the test dataset, through *detect.py*, and *val.py*. The small-scaled model has resulted in 48.5% precision rate, 21.8% recall rate, and a 22.6% mean average precision. The medium-scaled had a 30.3% precision rate, 26.4% recall rate, and a 25.6% mean average precision. Lastly, the large-scaled model had a 32.8% precision rate, 26.2% recall rate, and a 26.8% mean average precision.

<i>Models</i>	Precision	Recall	mAP
<i>Small</i>	48.5%	21.8%	22.6%
<i>Medium</i>	30.3%	26.4%	25.6%
<i>Large</i>	32.8%	26.2%	26.8%

As you can see, the small-scaled model scored better in precision but scored the lowest in the mean average precision rate. Note that mAP captures the tradeoff between precision and recall and maximizes the impact of both. Therefore, even with lower precision than the small-scaled model, the large-scaled model performed relatively better than all models as it has the highest mAP.

Furthermore, our trained model was able to detect motorcycle riders fairly, but it does not do a great job of classifying if they are wearing a proper helmet or not, or no helmet at all. As we can see from these [graphs](#), it does not do well on both the training and testing datasets. This could mean that the model is underfitted due to its poor performance on the training data and overfitted because it does not perform well on the evaluation data.

In the future, we recommend increasing the amount of training data examples, enhancing the video quality, and using fewer features for the model.

VI. Testing on a video from the previous assessment

Once the individual video frames (images) have been trained and tested, each video recording will now go through the same process with the final objective of validating the model's performance in real world scenarios (traffic surveillance). To do this, first, we need to install *opencv* to perform image processing on the video.

```
#install opencv-python
!pip install opencv-python
```

Then, we upload a specific video for testing. For this documentation, we will just be using one video.

```
#upload video file
from google.colab import files
video = files.upload()
```

[Browse...](#) XVR_ch9_main_20221004190008_20221004200000.mp4

XVR_ch9_main_20221004190008_20221004200000.mp4(video/mp4) - 77698947 bytes, last modified: n/a - 100% done

Saving XVR_ch9_main_20221004190008_20221004200000.mp4 to XVR_ch9_main_20221004190008_20221004200000.mp4

After uploading the video, we will then use the video to validate the model's performance.

```
!python detect.py --weights runs/train/yolov5m_results/weights/best.pt --
img 416 --conf 0.1 --source '/content/yolov5/
XVR_ch9_main_20221004190008_20221004200000.mp4'
```

detect: weights=['runs/train/yolov5m_results/weights/best.pt'], source=/content/yolov5/*.mp4, data=data/coco128.yaml, imgsz=[416, 416], conf_t
YOLOv5 🚀 v7.0-21-ga1b6e79 Python-3.8.15 torch-1.12.1+cu113 CUDA:0 (Tesla T4, 15110MiB)

Fusing layers...

Model summary: 212 layers, 20869098 parameters, 0 gradients, 47.9 GFLOPs

```
video 1/1 (1/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 14.7ms
video 1/1 (2/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 14.1ms
video 1/1 (3/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 14.0ms
video 1/1 (4/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 14.0ms
video 1/1 (5/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 14.2ms
video 1/1 (6/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 18.3ms
video 1/1 (7/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 15.4ms
video 1/1 (8/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 13.4ms
video 1/1 (9/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 13.3ms
video 1/1 (10/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 13.4ms
video 1/1 (11/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 13.3ms
video 1/1 (12/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 12.5ms
video 1/1 (13/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 13.0ms
video 1/1 (14/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 11.9ms
video 1/1 (15/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 14.3ms
video 1/1 (16/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 13.1ms
video 1/1 (17/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 15.2ms
video 1/1 (18/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 1 Rider, 12.3ms
```

```

video 1/1 (4480/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.4ms
video 1/1 (4481/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.5ms
video 1/1 (4482/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 14.8ms
video 1/1 (4483/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.9ms
video 1/1 (4484/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.9ms
video 1/1 (4485/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 19.3ms
video 1/1 (4486/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 11.7ms
video 1/1 (4487/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.9ms
video 1/1 (4488/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.4ms
video 1/1 (4489/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 11.7ms
video 1/1 (4490/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.7ms
video 1/1 (4491/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.9ms
video 1/1 (4492/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 11.8ms
video 1/1 (4493/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.4ms
video 1/1 (4494/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.8ms
video 1/1 (4495/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 10.8ms
video 1/1 (4496/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 11.4ms
video 1/1 (4497/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 12.3ms
video 1/1 (4498/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 12.1ms
video 1/1 (4499/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 12.1ms
video 1/1 (4500/4500) /content/yolov5/XVR_ch9_main_20221004190008_20221004200000.mp4: 256x416 (no detections), 11.7ms
Speed: 0.3ms pre-process, 12.0ms inference, 0.4ms NMS per image at shape (1, 3, 416, 416)
Results saved to runs/detect/exp5


```

Then, we download the results using the code snippet below:

```

files.download('/content/yolov5/runs/detect/exp5/XVR_ch9_main_20221004190008_20221004200000.mp4') #download the video after model detection

```



We repeated these steps until all models had the chance to be validated. If you would like to see all three video outputs for each model, the link is [here](#).

References

- Roboflow. (2020, June 15). *How to Train YOLO v5 on a Custom Dataset* [Video]. YouTube. <https://www.youtube.com/watch?v=MdF6x6ZmLAY>
- Solawetz, J. (2021, December 14). *How to Train A Custom Object Detection Model with YOLO v5*. Medium. <https://towardsdatascience.com/how-to-train-a-custom-object-detection-model-with-yolo-v5-917e9ce13208>