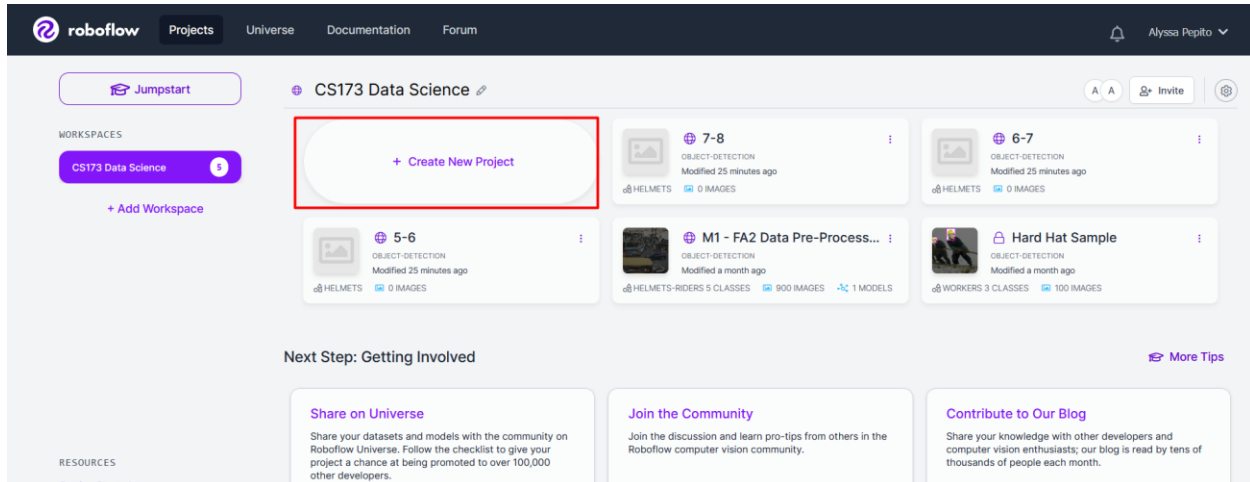
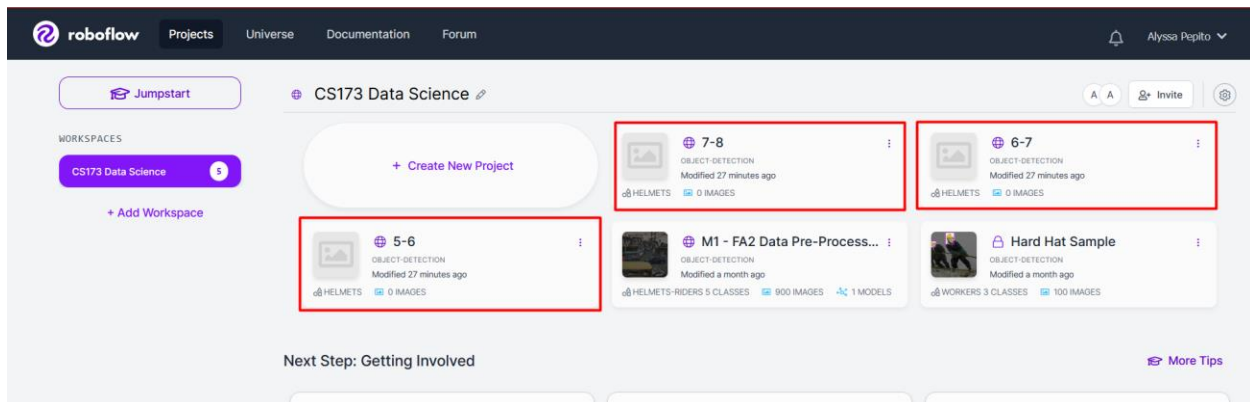


M2 - FA2 YOLOv7

I. Creating Three Separate Time Scenarios

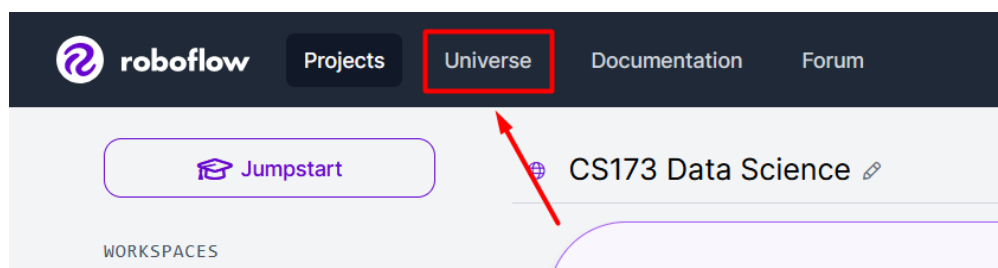


In Roboflow, we must first create three new projects to place our three-time scenarios. Click on “+ Create New Project”.




After filling up the details for the new project, repeat this step until you get three projects named specifically for the three-time scenarios (5-6, 6-7, and 7-8).

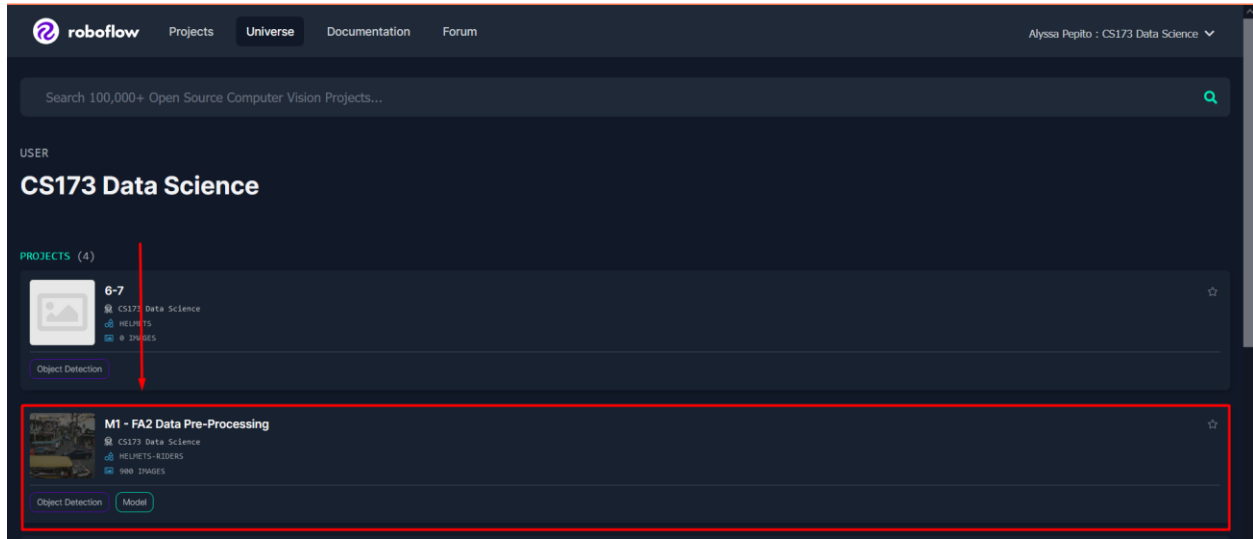
II. Cloning the Images



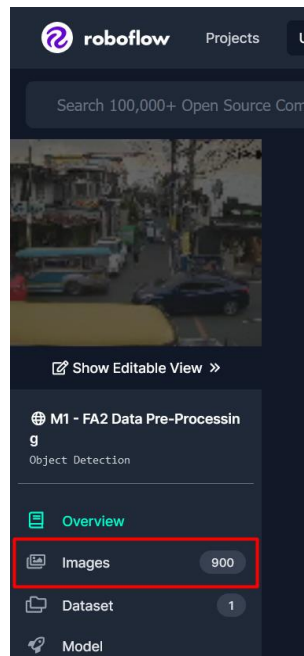
Next, on the top left corner of Roboflow, click on Universe.



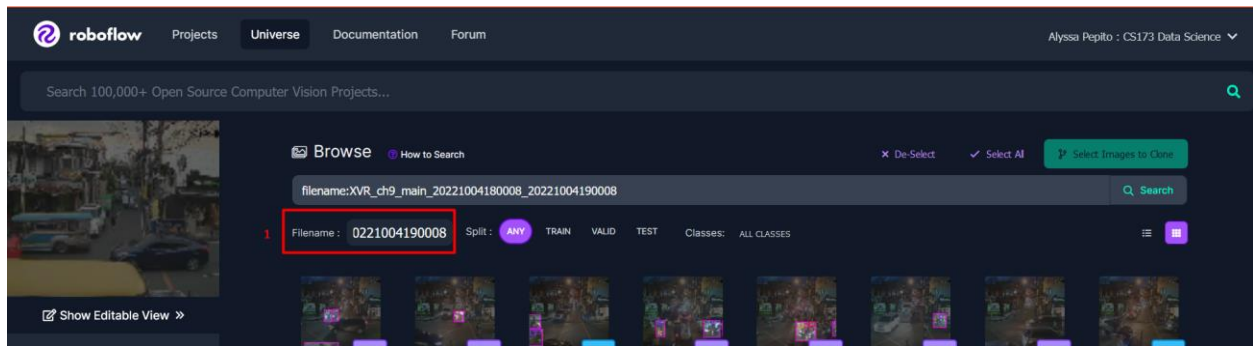
Roboflow Universe will open in a new tab. On the top right corner, click on the dropdown button and then click on “Workspace Profile”.



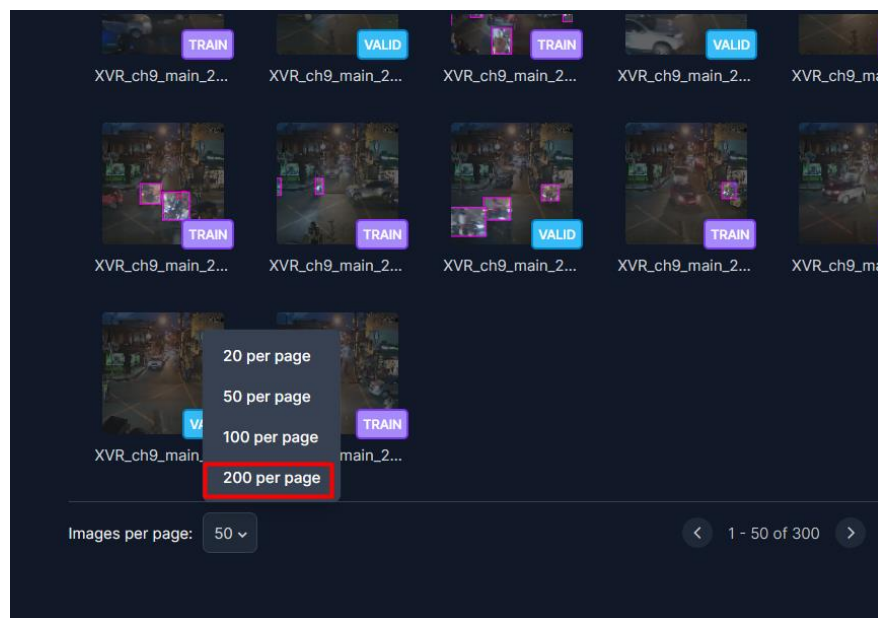
We select the old dataset we used from the previous activity “M1-FA2”.



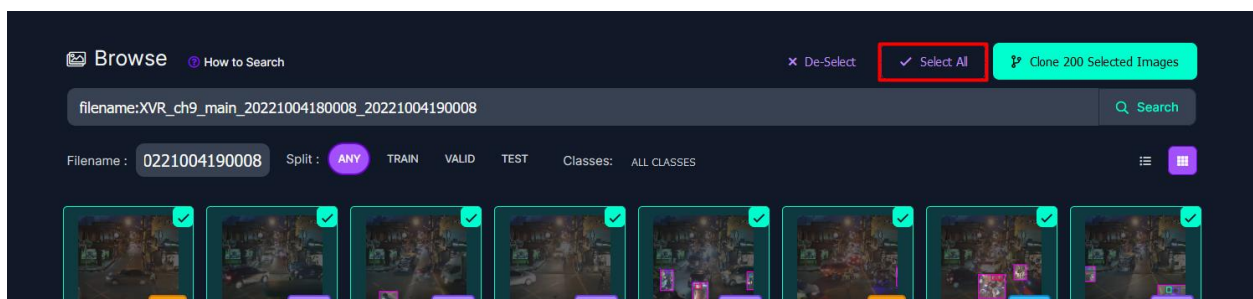
After the dataset has been loaded, we go to the left corner and click on “Images”.



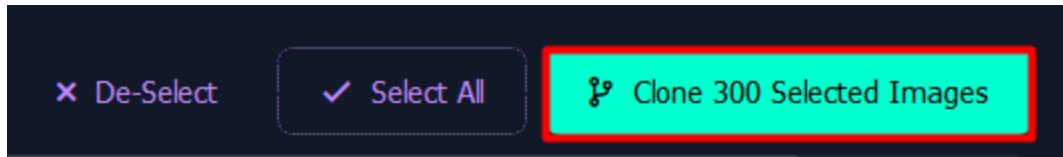
Next, we search for files that has the filename we want. As demonstrated above, we searched for “XVR_ch9_main_20221004190008_20221004200000” which are images from the time scenario 7-8 pm.



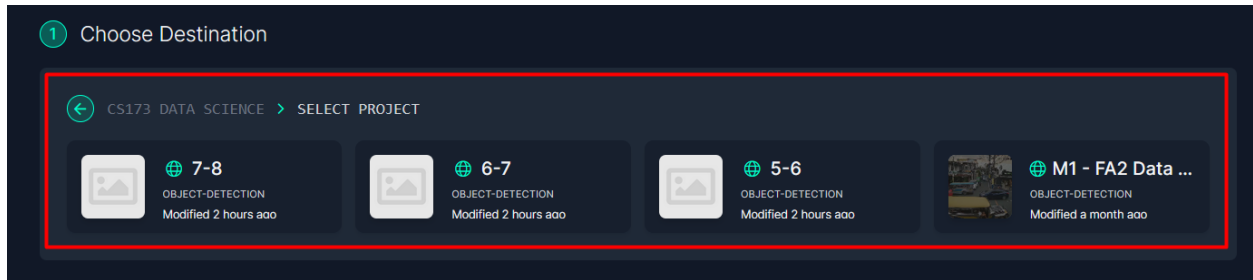
After the images have been loaded, we scrolled down to notice that there are a total of 300 images for that specific time scenario. To see more than 50 images, we clicked on the dropdown button next to the phrase “Images per page” and selected “200 per page”.



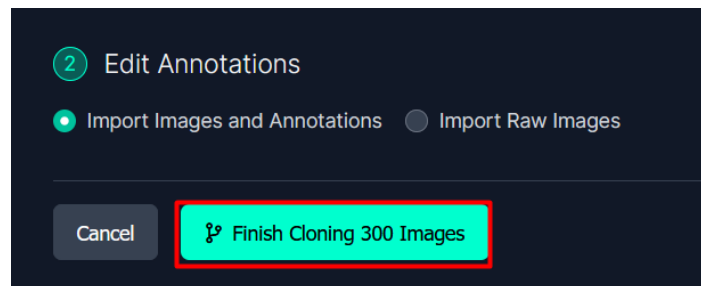
Once that there are 200 images in the first page, scroll up again and click on “Select all”. Go to the next page and click on “Select all” again until all 300 images are selected.



Next, click on “Clone 300 Selected Images” to start the cloning process.

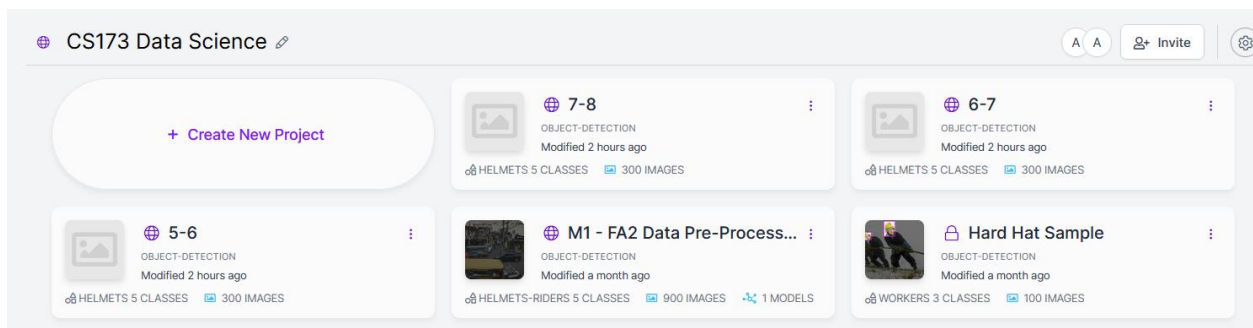


Scroll down to select which project to place the cloned images. In this case, it would be the 7-8 named project.




After selecting the project, we continue with the default option of “Import Images and Annotations” and click on “Finish Cloning 300 Images”. Repeat these [steps](#) until you have all images for every time scenario.

III. Generating New Versions of the Projects



We go back to Roboflow Projects to see all time scenario projects have all their images in place. We will then generate new versions for them. To start, we click on any one of them.



Inside the project, click on “+ Generate New Version” to get started.

Generating New Version		
Prepare your images and data for training by compiling them into a version. Experiment with different configurations to achieve better training results.		
✓ Source Images	Images: 300 Classes: 5 Unannotated: 0	
✓ Train/Test Split	Training Set: 209 images Validation Set: 62 images Testing Set: 29 images	
✓ Preprocessing	Auto-Orient: Applied Resize: Stretch to 640×640	
✓ Augmentation	Flip: Horizontal Crop: 0% Minimum Zoom, 20% Maximum Zoom Shear: ±15° Horizontal, ±15° Vertical Grayscale: Apply to 25% of images Hue: Between -25° and +25° Saturation: Between -25% and +25% Cutout: 3 boxes with 10% size each Mosaic: Applied Bounding Box: Crop: 0% Minimum Zoom, 20% Maximum Zoom Bounding Box: Brightness: Between -35% and +35%	

As seen above, we implemented auto-orient and resize for our preprocessing. For our augmentation, we applied flip, crop, shear, grayscale, hue, saturation, cutout, mosaic, bounding box: crop, and bounding box: brightness. Take note that one can experiment with their preprocessing and augmentation options. It is not necessary to follow what we have done here.

5

Generate

Review your selections and select a version size to create a moment-in-time snapshot of your dataset with the applied transformations.

Larger versions take longer to train but often result in better model performance. [See how this is calculated >](#)

Maximum Version Size

718 images (3x) ▼

Generate

Next, we continue with the default option of 3x for version size and click on Generate.

5-6 Image Dataset

+ Generate New Version

2023-01-06 2:37am

Version 1 Generated Jan 6, 2023

Export

Edit ⋮

VERSIONS

After generating, the dataset is ready to be used for training a model. To do that, we must click on “Export”.

Export



Format

1

YOLO v7 PyTorch ▼

TXT annotations and YAML config used with [YOLOv7](#).

☐ download zip to computer ☒ show download code 2

☒ Also train a model for [Label Assist](#) with [Roboflow Train](#).

Cancel

3

Continue

As seen above, make sure that the format selected should be YOLO v7 PyTorch. Select show download code and click on continue.

Your Download Code

Jupyter

Terminal

Raw URL

Paste this snippet into [a notebook from our model library](#) >> to download and unzip [your dataset](#) >>:

```
!pip install roboflow

from roboflow import Roboflow
rf = Roboflow(api_key="")
project = rf.workspace("").project("5-6")
dataset = project.version(1).download("yolov7")
```

Warning: Do not share this snippet beyond your team, it contains a private key that is tied to your Roboflow account. Acceptable use policy applies.

Done

The download code will then be generated, make sure to save this snippet for later. Repeat these [steps](#) for every time scenario project. By the end of this, you would have a total of three code snippets, each one from the three-time scenarios.

IV. Libraries and Dependencies Installation

Unlike the assessments we accomplished previously, this time we will be creating a YOLOv7 model. First, we cloned and downloaded the YOLOv7 repository and installed its dependencies. As shown in the code sample below, we also installed the Roboflow-specific library because the dataset for this model will originate from Roboflow.

```
# Download YOLOv7 repository and install requirements
!git clone https://github.com/WongKinYiu/yolov7
%cd yolov7
!pip install -r requirements.txt
%pip install -q roboflow # install roboflow

import torch
import os
from IPython.display import Image, clear_output # to display images

print(f"Setup complete. Using torch {torch.__version__} ({torch.cuda.get_device_properties(0).name if torch.cuda.is_available() else 'CPU'})")
```

Once the setup above is completed, we imported Roboflow by applying the following code snippet:

```
from roboflow import Roboflow #import roboflow for our datasets
rf = Roboflow(model_format="yolov7", notebook="ultralytics")
```

We then set up the environment.

```
# set up environment
os.environ["DATASET_DIRECTORY"] = "/content/datasets"
```

Choose any of the code snippets we have [copied previously](#) from the three-time scenarios for now, as we must repeat this with the other time scenarios later. Then, paste the snippet below the environment set-up and run it. The Roboflow dataset will be downloading.

```
from roboflow import Roboflow
rf = Roboflow(api_key="RU6eGLCNP3Q4UwQqKWha")
project = rf.workspace("cs173-data-science").project("5-6")
dataset = project.version(1).download("yolov7")
```

After the Roboflow dataset has download, we will then download the required weights for our training.

```
# download COCO starting checkpoint
%cd /content/yolov7
!wget https://github.com/WongKinYiu/yolov7/releases/download/v0.1/yolov7_training.pt
```

V. Training the YOLOv7 model

After installing all required files and dependencies, we used train.py to train the model. The following hyperparameters were chosen for this model:

1. Image size = 640 pixels
2. Batch size = 16
3. Epochs = 50
4. YOLOv7 Version = YOLOv7

The following code snippet was executed to apply the ideal hyperparameters specified above, producing the results from the associated figures:

```
!python train.py --img 640 --batch 16 --epochs 50 --
data {dataset.location}/data.yaml --weights 'yolov7_training.pt' --cache
```



```
!python train.py --img 640 --batch 16 --epochs 50 --data {dataset.location}/data.yaml --weights 'yolov7_training.pt' --cache
```

YOLOR v0.1-121-g2fdc7f1 torch 1.13.0+cu116 CUDA:0 (Tesla T4, 15109.75MB)

Namespace(adam=False, artifact_alias='latest', batch_size=16, bbox_interval=-1, bucket='', cache_images=True, cfg='', data='/content/datasets/
tensorboard: Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>
hyperparameters: lr=0.01, lrf=0.1, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=0.05,
wandb: Install Weights & Biases for YOLOR logging with 'pip install wandb' (recommended)
Overriding model.yaml nc=80 with nc=5

	from	n	params	module	arguments
0	-1	1	928	models.common.Conv	[3, 32, 3, 1]
1	-1	1	18560	models.common.Conv	[32, 64, 3, 2]
2	-1	1	36992	models.common.Conv	[64, 64, 3, 1]
3	-1	1	73984	models.common.Conv	[64, 128, 3, 2]
4	-1	1	8320	models.common.Conv	[128, 64, 1, 1]
5	-2	1	8320	models.common.Conv	[128, 64, 1, 1]
6	-1	1	36992	models.common.Conv	[64, 64, 3, 1]
7	-1	1	36992	models.common.Conv	[64, 64, 3, 1]
8	-1	1	36992	models.common.Conv	[64, 64, 3, 1]
9	-1	1	36992	models.common.Conv	[64, 64, 3, 1]
10	[-1, -3, -5, -6]	1	0	models.common.Concat	[1]
11	-1	1	66048	models.common.Conv	[256, 256, 1, 1]
12	-1	1	0	models.common.MP	[1]
13	-1	1	33024	models.common.Conv	[256, 128, 1, 1]
14	-3	1	33024	models.common.Conv	[256, 128, 1, 1]
15	-1	1	147712	models.common.Conv	[128, 128, 3, 2]

Transferred 557/566 items from yolov7_training.pt

Scaled weight_decay = 0.0005

Optimizer groups: 95 .bias, 95 conv.weight, 98 other

train: Scanning '/content/datasets/5-6-1/train/labels' images and labels... 615 found, 0 missing, 2 empty, 0 corrupted: 100% 615/615 [00:00<00

train: New cache created: /content/datasets/5-6-1/train/labels.cache

train: Caching images (0.8GB): 100% 615/615 [00:02<00:00, 238.19it/s]

val: Scanning '/content/datasets/5-6-1/valid/labels' images and labels... 62 found, 0 missing, 0 empty, 0 corrupted: 100% 62/62 [00:00<00:00,

val: New cache created: /content/datasets/5-6-1/valid/labels.cache

val: Caching images (0.1GB): 100% 62/62 [00:00<00:00, 176.90it/s]

autoanchor: Analyzing anchors... anchors/target = 4.86, Best Possible Recall (BPR) = 0.9990

Image sizes 640 train, 640 test

Using 2 dataloader workers

Logging results to runs/train/exp

Starting training for 50 epochs...

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
0/49	8G	0.08582	0.0261	0.0256	0.1375	156	640: 100% 39/39 [01:05<00:00, 1.67s/it]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 2/2 [00:17<00:00, 8.62s/it]
	all	62	436	0.00184	0.00622	0.000483	7.35e-05

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
1/49	10.8G	0.07599	0.02422	0.01792	0.1181	105	640: 100% 39/39 [00:32<00:00, 1.18it/s]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.55it/s]
	all	62	436	0.214	0.0653	0.00912	0.00169

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
46/49	10.8G	0.0264	0.01843	0.004139	0.04897	159	640: 100% 39/39 [00:33<00:00, 1.18it/s]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.55it/s]
	all	62	436	0.556	0.355	0.31	0.149

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
47/49	10.8G	0.0264	0.0187	0.004057	0.04915	91	640: 100% 39/39 [00:34<00:00, 1.15it/s]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.53it/s]
	all	62	436	0.653	0.361	0.314	0.147

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
48/49	10.8G	0.02592	0.01752	0.004138	0.04758	149	640: 100% 39/39 [00:33<00:00, 1.18it/s]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.54it/s]
	all	62	436	0.717	0.335	0.314	0.153

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
49/49	10.8G	0.02603	0.01747	0.004216	0.04772	72	640: 100% 39/39 [00:32<00:00, 1.19it/s]
	Class	Images	Labels		P	R	mAP@.5 mAP@.5:.95: 100% 2/2 [00:02<00:00, 1.10s/it]
	all	62	436	0.549	0.38	0.322	0.155
	Full-Faced	62	109	0.406	0.596	0.423	0.114
	Half-Faced	62	86	0.256	0.442	0.232	0.0865
	Invalid	62	22	0.328	0.0455	0.0486	0.0167
	Not Wearing Helmet	62	26	1	0	0.0527	0.0152
	Rider	62	193	0.754	0.819	0.855	0.543

50 epochs completed in 0.520 hours.

Optimizer stripped from runs/train/exp/weights/last.pt, 74.8MB

Optimizer stripped from runs/train/exp/weights/best.pt, 74.8MB

Since the above was resulted from the 5:00 – 6:00 pm time scenario, we would need to change this snippet to another one from the other time scenarios. Then, execute train.py with the same hyperparameters. The following are the results from the 6:00 – 7:00 pm time scenario.

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
46/49	10.8G	0.02333	0.01155	0.00246	0.03734	39	640: 100% 41/41 [00:32<00:00, 1.26it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.91it/s]
	all	56	158	0.767	0.323	0.346	0.147
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
47/49	10.8G	0.02352	0.01185	0.002305	0.03767	47	640: 100% 41/41 [00:32<00:00, 1.26it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.92it/s]
	all	56	158	0.836	0.288	0.348	0.15
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
48/49	10.8G	0.02275	0.01127	0.002099	0.03613	25	640: 100% 41/41 [00:33<00:00, 1.22it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.96it/s]
	all	56	158	0.81	0.323	0.364	0.159
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
49/49	10.8G	0.02313	0.01139	0.002152	0.03667	50	640: 100% 41/41 [00:32<00:00, 1.26it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.24it/s]
	all	56	158	0.828	0.297	0.365	0.159
	Full-Faced	56	32	0.515	0.332	0.4	0.13
	Half-Faced	56	18	0.8	0.444	0.511	0.147
	Invalid	56	1	1	0	0.0138	0.00553
Not Wearing Helmet	56	1	1	0	0.111	0.0553	
Rider	56	106	0.825	0.708	0.793	0.459	

50 epochs completed in 0.512 hours.

Optimizer stripped from runs/train/exp/weights/last.pt, 74.8MB
Optimizer stripped from runs/train/exp/weights/best.pt, 74.8MB

Lastly, these are the results from the 7:00 – 8:00 pm time scenario.

Epoch	gpu_mem	box	obj	cls	total	labels	img_size
43/49	10.9G	0.02788	0.01207	0.003985	0.04394	28	640: 100% 40/40 [00:33<00:00, 1.21it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.70it/s]
	all	62	208	0.64	0.303	0.245	0.12
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
44/49	10.9G	0.02654	0.01217	0.003796	0.04251	20	640: 100% 40/40 [00:31<00:00, 1.25it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.71it/s]
	all	62	208	0.665	0.277	0.248	0.12
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
45/49	10.9G	0.02685	0.01213	0.003628	0.0426	27	640: 100% 40/40 [00:32<00:00, 1.25it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.71it/s]
	all	62	208	0.64	0.275	0.231	0.117
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
46/49	10.9G	0.0266	0.0119	0.003459	0.04196	22	640: 100% 40/40 [00:31<00:00, 1.25it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.73it/s]
	all	62	208	0.64	0.321	0.253	0.121
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
47/49	10.9G	0.02702	0.01205	0.003789	0.04285	25	640: 100% 40/40 [00:32<00:00, 1.25it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.71it/s]
	all	62	208	0.661	0.292	0.254	0.123
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
48/49	10.9G	0.02689	0.01148	0.003695	0.04206	51	640: 100% 40/40 [00:31<00:00, 1.26it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.71it/s]
	all	62	208	0.663	0.29	0.254	0.121
Epoch	gpu_mem	box	obj	cls	total	labels	img_size
49/49	10.9G	0.02567	0.01163	0.003542	0.04085	3	640: 100% 40/40 [00:31<00:00, 1.25it/s]
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95: 100% 2/2 [00:01<00:00, 1.08it/s]
	all	62	208	0.638	0.311	0.245	0.117
	Full-Faced	62	37	0.174	0.405	0.211	0.0476
	Half-Faced	62	54	0.338	0.368	0.232	0.0512
	Invalid	62	4	1	0	0.00203	0.000862
Not Wearing Helmet	62	2	1	0	0	0	0
Rider	62	111	0.681	0.784	0.78	0.485	

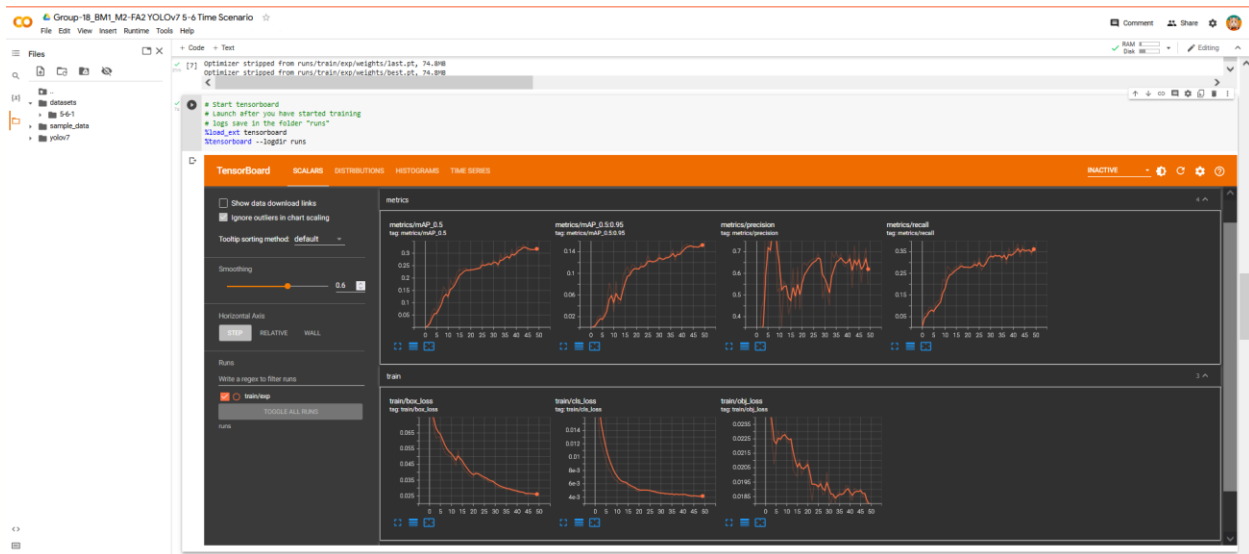
50 epochs completed in 0.503 hours.

Optimizer stripped from runs/train/exp/weights/last.pt, 74.8MB
Optimizer stripped from runs/train/exp/weights/best.pt, 74.8MB

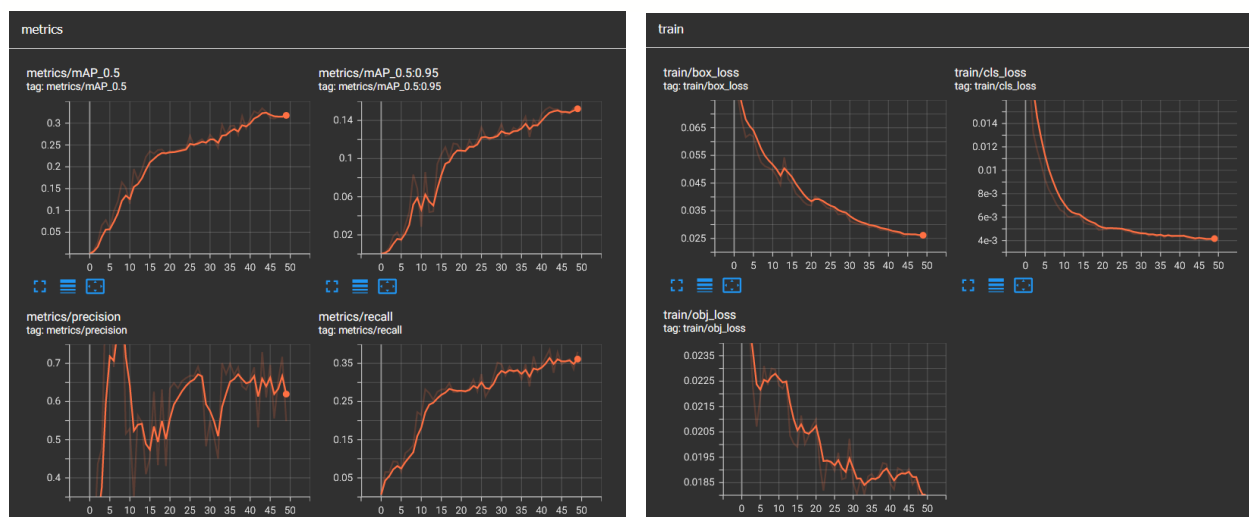
VI. TensorBoard


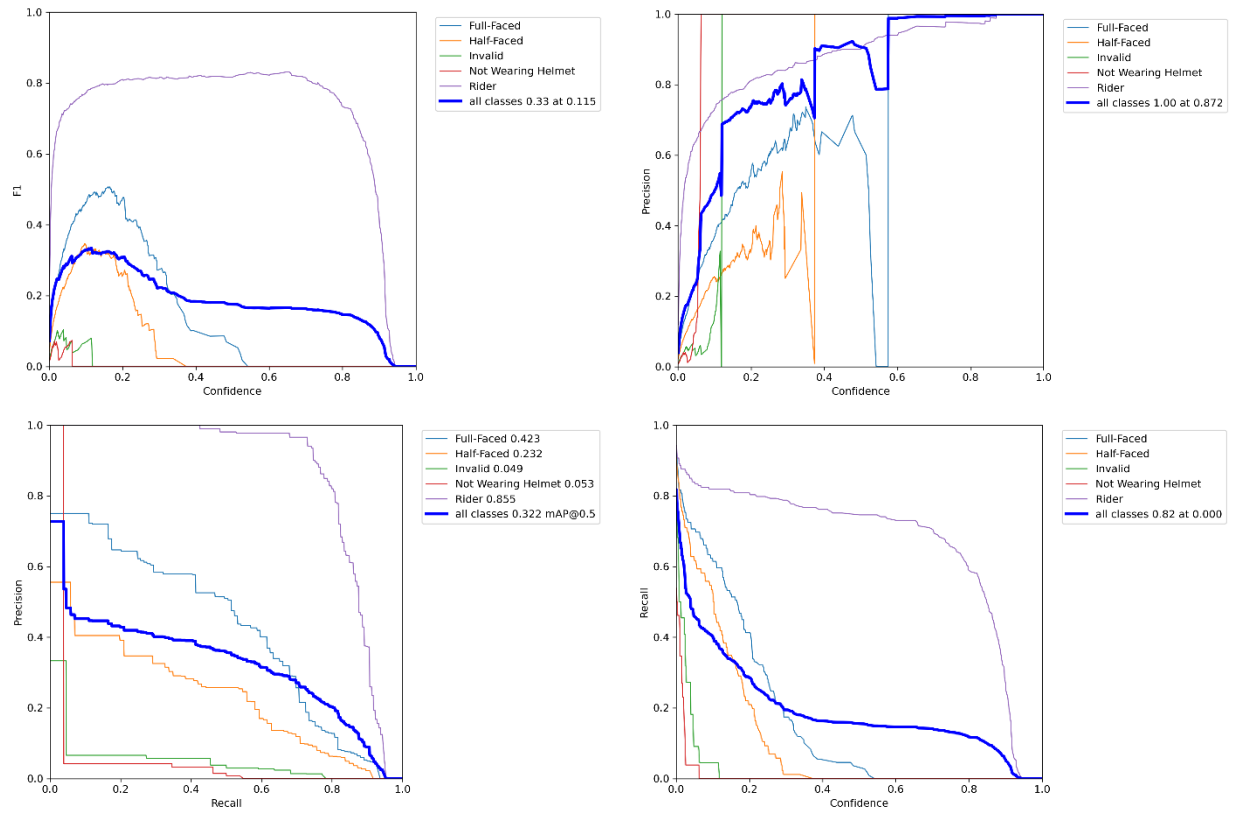
TensorBoard was once again employed to visualize the training process' outcomes using scalar measurements, pictures, graphs, and time series. Scalar measures like mAP (mean average precision), precision, and recall are just a few examples. One can view the model's performance through graphs that show the metrics, train, and the model's overall outcomes under. The time series displays each runtime that the model underwent during training. The TensorBoard was started by running the following line of code:

```
%load_ext tensorboard
%tensorboard --logdir runs
```










The graphical results for the 5:00 – 6:00 pm time scenario are seen in the figures below:



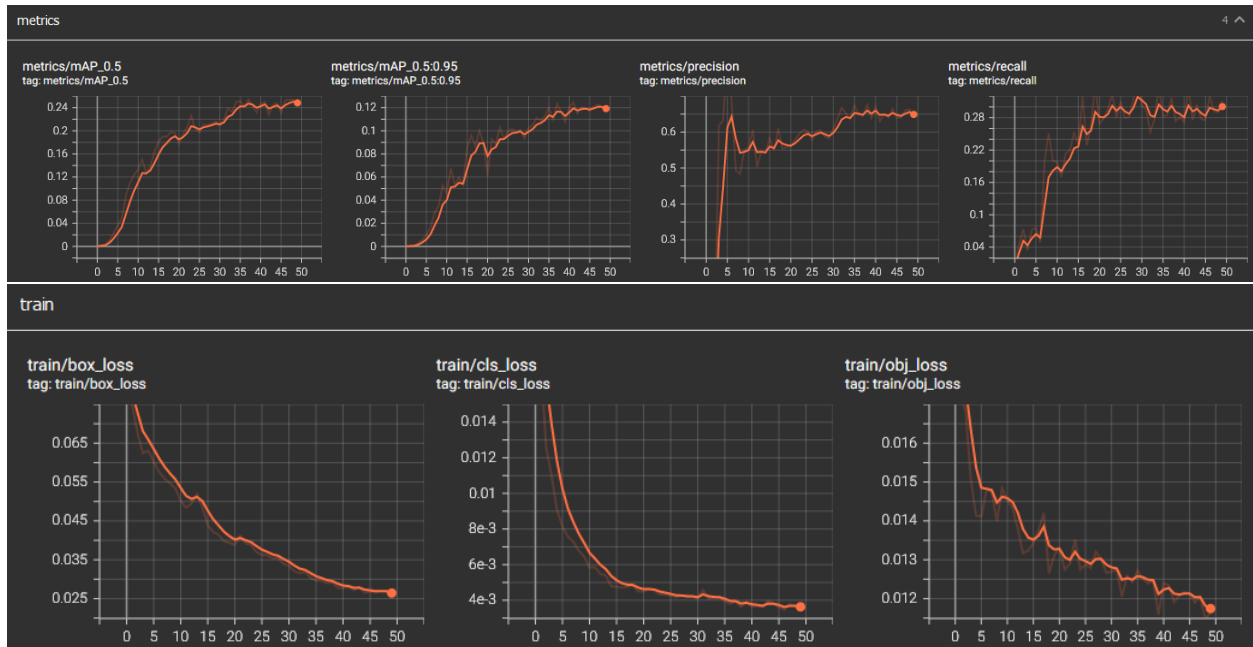







The results for the 6:00 – 7:00 pm time scenario are seen below:





Finally, the results for the 7:00 - 8:00 pm time scenario are in the following figures:





VII. Testing the YOLOv7 model

The testing dataset will be used to test the model after it has been trained using the training dataset. This dataset's goal is to evaluate the model's performance indicators and detection abilities. In order to achieve this goal, *detect.py* was used to identify objects and predict the placement of the bounding boxes from the test photos.

```
!python detect.py --weights runs/train/exp/weights/best.pt --img 640 --
conf 0.1 --source {dataset.location}/test/images
```

```
Namespace(agnostic_nms=False, augment=False, classes=None, conf_thres=0.1, device='', exist_ok=False, img_size=640, iou_thres=0.45, name='exp', no_trace=False,
YOLOv7 v0.1-121-g2fcd7f1 torch 1.13.0+cu116 CUDA:0 (Tesla T4, 15109.75MB)

Fusing layers...
RepConv.fuse_repvgg_block
RepConv.fuse_repvgg_block
RepConv.fuse_repvgg_block
IDetect.fuse
/usr/local/lib/python3.8/dist-packages/torch/functional.py:504: UserWarning: torch.meshgrid: in an upcoming release, it will be required to pass the index
return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
Model Summary: 314 layers, 36503348 parameters, 6194944 gradients, 103.2 GFLOPS
Convert model to Traced-model...
traced_script_module saved!
model is traced!

5 Full-Faceds, 3 Riders, Done. (23.1ms) Inference, (1.4ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-100.jpg.rf.daa98d56597398c66bb44194dc0bbc21.jpg
5 Full-Faceds, 2 Riders, Done. (23.1ms) Inference, (1.0ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-108.jpg.rf.ad2191a305b2f2457c8874046e14ba71.jpg
6 Full-Faceds, 2 Half-Faceds, 5 Riders, Done. (23.0ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-112.jpg.rf.eac3bb893906371842333a806067ca73.jpg
1 Full-Faced, 1 Half-Faced, 3 Riders, Done. (23.1ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-121.jpg.rf.5843c6c85a690eb6e2ef1e2c22b80136.jpg
2 Full-Faceds, 5 Half-Faceds, 5 Riders, Done. (23.0ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-127.jpg.rf.b2d9966ebb02b74989ca35a2d22f032b.jpg
1 Half-Faced, 2 Riders, Done. (23.1ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-144.jpg.rf.27949ead6bcdcc806db94e4d771ba949.jpg

2 Full-Faceds, 2 Half-Faceds, 3 Riders, Done. (12.9ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-254.jpg.rf.9c72b8cf1576c5050d150b39f1d697dd.jpg
5 Full-Faceds, 6 Half-Faceds, 5 Riders, Done. (12.9ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-256.jpg.rf.04eadbfa3474776edef5cd3142cbc080.jpg
4 Full-Faceds, 2 Half-Faceds, 5 Riders, Done. (12.9ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-259.jpg.rf.fb659b404c1eb9b815d4f5c74431160c.jpg
4 Full-Faceds, 6 Half-Faceds, 8 Riders, Done. (13.0ms) Inference, (1.3ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-264.jpg.rf.76112cb5a1324ab66f0b9327d2a78cb1.jpg
9 Half-Faceds, 7 Riders, Done. (12.7ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-284.jpg.rf.fa48a91ec68a629599811c07d530893b.jpg
5 Half-Faceds, 7 Riders, Done. (12.7ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-291.jpg.rf.6c01e293d15a444ba292fb3ba21ed7a73.jpg
1 Half-Faced, 4 Riders, Done. (12.6ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-293.jpg.rf.200ac70a660e479c97011471499a1401.jpg
2 Half-Faceds, 2 Riders, Done. (12.1ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-35.jpg.rf.5595b60e35b410a207212ec9b7253f42.jpg
2 Full-Faceds, 3 Half-Faceds, 5 Riders, Done. (12.1ms) Inference, (0.9ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-45.jpg.rf.d4be1e46a55b6d49f239272f006b49d0.jpg
1 Half-Faced, 3 Riders, Done. (12.2ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-54.jpg.rf.f0399238194c4f79a01b749a21b9aff7.jpg
1 Full-Faced, 1 Half-Faced, 1 Rider, Done. (12.1ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-62.jpg.rf.c8f3b567ad5bc6d99f9a6dcc8fdda8e1.jpg
3 Riders, Done. (12.1ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-72.jpg.rf.d173c3ca1d4811550092dda660ed6fd4.jpg
6 Full-Faceds, 1 Half-Faced, 5 Riders, Done. (12.2ms) Inference, (0.8ms) NMS
The image with the result is saved in: runs/detect/exp/XVR_ch9_main_20221004170008_20221004180008_mp4-90.jpg.rf.82c054c354a9cb8a24431454e60a9cd0.jpg
Done. (1.074s)
```

To display all test images with the bounding boxes placed on their predictions, the following code was used:

```
import glob
from IPython.display import Image, display

for imageName in glob.glob('/content/yolov7/runs/detect/exp/*.jpg'): #assuming JPG
    display(Image(filename=imageName))
```



```
print("\n")
```

The figures below are a few examples of the inferred (with bounding boxes) test images.



The table output seen in the image below was then created by running the following code. The code snippet will generate a summary of the model's values through a table, as its classes are divided into Full-Faced, Half-Faced, Invalid, Not Wearing Helmet, and Rider and classified according to the number of images, instances, precision (P), recall (R), mAP50, and mAP50-95. The results will be saved as a file named "best.pt."

```
!python test.py --weights './runs/train/exp/weights/best.pt' --
data {dataset.location}/data.yaml --img 640
```

Output for the 5:00 – 6:00 pm time scenario:

Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95:
all	62	436	0.537	0.382	0.317	0.153
Full-Faced	62	109	0.375	0.596	0.413	0.111
Half-Faced	62	86	0.242	0.449	0.225	0.0825
Invalid	62	22	0.328	0.0448	0.0473	0.0162
Not Wearing Helmet	62	26	1	0	0.0494	0.0144
Rider	62	193	0.74	0.819	0.852	0.541

Output for the 6:00 – 7:00 pm time scenario:

Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95:
all	56	158	0.897	0.298	0.356	0.168
Full-Faced	56	32	0.611	0.344	0.402	0.172
Half-Faced	56	18	1	0.441	0.536	0.2
Invalid	56	1	1	0	0.0105	0.00629
Not Wearing Helmet	56	1	1	0	0.0356	0.0213
Rider	56	106	0.872	0.708	0.797	0.442

Output for the 7:00 – 8:00 pm time scenario:

Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95:
all	62	208	0.649	0.274	0.244	0.123
Full-Faced	62	37	0.174	0.297	0.179	0.0494
Half-Faced	62	54	0.374	0.278	0.248	0.0703
Invalid	62	4	1	0	0.000769	0.000385
Not Wearing Helmet	62	2	1	0	0	0
Rider	62	111	0.697	0.793	0.791	0.495

VIII. Ground Truth of Dataset and Confusion Matrix

We would be using the same method for manually counting all classes within the test dataset. Using excel to accelerate the process, the following tables were created:


5:00PM - 6:00PM		
Total		
Rider		99
Full-Faced		53
Half-Faced		41
Invalid		14
Not Wearing Helmet		6

6:00PM - 7:00PM		
Total		
Rider		69
Full-Faced		24
Half-Faced		13
Invalid		0
Not Wearing Helmet		1


7:00PM-8:00PM		
Total		
Rider		66
Full-Faced		35
Half-Faced		23
Invalid		2
Not Wearing Helmet		0

After testing the YOLOv7 model, we were able to produce these confusion matrices.


5:00 – 6:00 PM



6:00 – 7:00 PM



7:00 – 8:00 PM



IX. Collab Notebooks

Below are the links to our collab notebooks:

1. 5:00 – 6:00 pm scenario

<https://colab.research.google.com/drive/1rtCEkPe7VSbuuEtunktTeXTIKOmLNiao?usp=sharing>

2. 6:00 – 7:00 pm scenario

<https://colab.research.google.com/drive/1tYhQjceTyuyJ1Mg4mHfg-Otl-gLM1UdW?usp=sharing>

3. 7:00 – 8:00 pm scenario

<https://colab.research.google.com/drive/1S7k35CjCrOshuz2DWnFThMaX7edaEBOc?usp=sharing>

References

Dwyer, B. (2022, December 29). *How to Train YOLOv7 on a Custom Dataset*. Roboflow Blog.

<https://blog.roboflow.com/yolov7-custom-dataset-training-tutorial/>

Skelton, J. (2022, August 17). *How to train and use a custom YOLOv7 model*. Paperspace Blog.

<https://blog.paperspace.com/yolov7/>